

Collateral Damage in Graph-based Memory Forgetting: Quantification, Exploitation, and Defense

Anonymous Authors¹

Abstract

Agent memory systems increasingly adopt entity graphs, and graph-based invalidation propagation—cascading updates through entity connections—is a plausible extension for consistency maintenance. We show that this step is dangerous without safeguards: **structural co-occurrence does not imply dependency**, and naive BFS propagation over co-occurrence graphs applies nonzero weight decay to 69–79% of valid memories across 6K–64K turns, while the two highest named-entity cross-task hubs cause –10.0pp and –11.9pp F1 drops on unrelated retrieval. Under adversarial hub-targeting, five injected facts apply nonzero weight decay to 57.3% of valid memories—weaponizing the system’s own maintenance mechanism. We propose **ATTR-AWARE Selective Propagation**, an attribute-keyed directional filter that restricts cascades to plausible dependency paths without LLM calls, achieving 78–100% damage reduction. Ablation indicates that directional filtering accounts for the damage reduction in our controlled attack setting; in-degree dampening affects only severity. Our initial findings on MemoryAgentBench suggest that systems exploring graph-based forgetting over co-occurrence structures should strongly consider dependency-aware propagation mechanisms from the outset to mitigate collateral damage.

1. Introduction

AI agents with persistent memory are moving from prototypes to production. Mem0 (Chhikara et al., 2025) and other memory-augmented agent frameworks have proliferated across open-source and proprietary ecosystems (Yang et al., 2026). These systems increasingly maintain entity

graphs to organize stored facts (Chhikara et al., 2025; Jiang et al., 2026; Rasmussen et al., 2025).

As memory scales to tens of thousands of facts, maintaining consistency becomes critical—yet multi-hop conflict resolution accuracy remains below 6% on standard benchmarks (Hu et al., 2026). The gap is structural: current systems handle invalidation *individually*, leaving dependent facts silently stale. This creates concrete pressure toward propagation-style maintenance. Community issue reports suggest that per-fact-only invalidation can lead to accumulating inconsistencies over time.¹ Graph-based propagation—cascading updates through entity connections—is a plausible extension of current graph-augmented memory systems. Zep (Rasmussen et al., 2025) maintains temporal validity windows on graph edges; MAGMA (Jiang et al., 2026) constructs explicit causal sub-graphs; Mem0 (Chhikara et al., 2025) already traverses entity graphs during conflict detection. Existing traversal infrastructure may lower the implementation barrier for propagation-style invalidation. Each system implements prerequisites for propagation without taking the final step.

We implement this step and reveal a fundamental tension: **co-occurrence graphs conflate correlation with causation**. Facts sharing an entity are linked identically regardless of dependency—“the president of Russia” connects diplomatic and economic facts through the same edges. This is not an artifact of naive graph construction: even typed-relation graphs contain co-occurrence as an unavoidable substructure whenever multiple facts reference the same entity. BFS propagation treats all such connections identically, cascading invalidation through semantically meaningless edges.

Because graph-augmented memory systems already maintain the prerequisites for traversal-based maintenance, unfiltered propagation is a plausible next design step rather than a deployed standard. Our study evaluates this step before deployment: on MemoryAgentBench-derived co-occurrence graphs, it applies nonzero weight decay to 69–79% of valid memories and motivates dependency-aware

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

¹Mem0 GitHub Issues #4573, #4863, #3841 (2025) describe stale entity links and contradictory facts as recurring community complaints.

safeguards. This distinction is well-understood in transportation network resilience (Chen et al., 2007) but has not been applied to AI memory systems. We contribute:

1. **Quantified risk:** BFS propagation over co-occurrence graphs applies nonzero weight decay to 69–79% of valid memories across 6K–64K scales; named-entity hub triggers cause −10.0pp and −11.9pp F1 drops on unrelated retrieval (Section 4).
2. **Adversarial amplification:** Five injected facts exploit hub topology to apply nonzero weight decay to 57.3% of valid memories; the attack requires gray-box access (hub entity names) and weaponizes the system’s own forgetting mechanism (Section 5).
3. **Practical defense:** ATTR-AWARE propagation, a hyperparameter-free attribute-keyed directional filter, reduces collateral damage by 78–100% without LLM calls. Ablation shows that directional filtering accounts for the damage reduction in our controlled attack setting (Section 6).

We characterize these failure modes *before* unfiltered graph propagation becomes a common maintenance strategy—providing initial quantified evidence and a working defense for systems that explore propagation beyond per-fact invalidation.

2. Background

Graph-augmented memory systems. Modern AI memory systems differ in how they encode entity relationships. Table 1 summarizes existing architectures along the dimension most relevant to propagation: edge semantics and invalidation strategy. In their published descriptions, graphs serve *retrieval* and consolidation but none specifies a general-purpose *invalidation propagation* mechanism over entity co-occurrence edges.

System	Graph	Edges	Invalidation
Mem0	KG	Typed	Per-fact
Zep	Temporal KG	Bi-temporal	Per-edge
MAGMA	4 sub-graphs	Causal+Entity	Per-graph
A-Mem	Linked notes	Associative	Per-note

Table 1. Graph-augmented memory systems (Chhikara et al., 2025; Rasmussen et al., 2025; Jiang et al., 2026; Xu et al., 2025). We found no explicit invalidation propagation in their published descriptions.

The forgetting gap. MemoryAgentBench (Hu et al., 2026) evaluates multi-hop conflict resolution where systems must update or discard outdated facts under contradictory

evidence. No tested system exceeds 6% accuracy. MemoryArena (He et al., 2026) shows agents drop to 40–60% on session-interdependent tasks. The failure mode is consistent: systems cannot reason about *dependencies between facts*.

Belief revision perspective. Kumiho (Park, 2026) formalizes this problem through AGM belief revision (Alchourrón et al., 1985): typed dependency edges (e.g., *Depends_On*, *Derived_From*) ensure that only causally related beliefs are affected during revision. This provides axiomatic guarantees—but requires typed dependency edges at graph construction time, which Mem0, Zep, and MAGMA do not currently provide. Our work addresses the complementary question: *what happens when propagation is attempted on existing untyped co-occurrence graphs?*

Memory as an attack surface. MINJA (Dong et al., 2025) achieves 98% injection success via query-only interaction; AgentPoison (Chen et al., 2024) poisons 80%+ of agent decisions by modifying <0.1% of stored entries. Xu et al. (2026) show that memory retrieval can steer agent control flow, while Devarangadi Sunil et al. (2026) evaluate poisoning defenses under realistic conditions with pre-existing legitimate memories. Xiong et al. (2025) empirically demonstrate that memory addition/deletion policies significantly impact long-term agent performance. MITRE ATLAS designates memory manipulation as technique AML.T0080 (MITRE Corporation, 2023). These attacks inject malicious *content*; our work reveals a structurally distinct threat where the system’s own *maintenance mechanism* becomes the attack vector.

Structural vulnerability from graph topology. The content-injection attacks above exploit *what* is stored; a complementary question is whether the graph *structure itself* creates vulnerability. Networks with heavy-tailed degree distributions are robust to random failures but catastrophically fragile to targeted hub removal (Albert et al., 2000; Cohen et al., 2000). Transportation network resilience provides a structural analogy: directional failure models separate correlated from causally dependent links (Chen et al., 2007), though the dynamics differ between physical and information networks. We test whether memory co-occurrence graphs exhibit this topology and whether hub-targeted operations produce the predicted damage.

3. The Propagation Trade-off

Before analyzing failure modes, we establish *why* propagation matters. We design 10 fact-change scenarios (personnel changes, schedule updates) with 5 causally related and 15 unrelated memories each. For instance, updating “the CEO of Acme is Alice” to “the CEO of Acme is Bob” should invalidate “Alice approves Acme milestones” (dependent)

but not “Acme was founded in 2010” (co-occurring but independent).

Mode	Forget Acc.	Collateral	Stale@worst
No-Propagation	84.0%	0.0%	40%
BFS	96.0%	47.5%	0%
ATTR-AWARE	94.0%	25.0%	8%

Table 2. Forgetting benefit vs. collateral cost (10 controlled scenarios). No-Propagation leaves 16% stale (worst 40%). BFS achieves highest accuracy but damages 47.5% of unrelated memories. ATTR-AWARE reduces collateral, though 25% residual indicates room for improvement.

Finding 1. Table 2 shows that No-Propagation leaves substantial stale content (16% average, 40% worst case). BFS eliminates staleness but damages 47.5% of unrelated memories—producing more damage than the staleness it eliminates. ATTR-AWARE achieves near-BFS forgetting accuracy (94% vs 96%) while reducing collateral damage, demonstrating that *if propagation is adopted*, dependency-aware filtering is essential.

The remainder of this paper characterizes the failure modes that make unfiltered propagation dangerous (Sections 4 and 5) and validates the defense (Section 6).

4. Collateral Damage at Scale

4.1. Setup

We run FactConsolidation memorization on MemoryAgentBench (Hu et al., 2026) under three modes: Baseline (no propagation), BFS (BFS through co-occurrence graph), and ATTR-AWARE (attribute-directed chains only). Infrastructure: Gemma-4-31B-it (LLM), BAAI/bge-m3 1024-dim (embedding), NetworkX graph, depth=2, decay_per_hop=0.5. No current system deploys BFS propagation as a standard invalidation policy; we implement it as a stress test of a plausible next step for graph-augmented memory systems that seek consistency beyond per-fact updates. Our results should therefore be read as proactive risk characterization, not as a claim about current production behavior.

We define five metrics: *damage*—any nonzero memory weight reduction ($w < 1.0$) from propagation; *severe damage*—weight reduced to retrieval-degrading levels ($0.1 \leq w < 0.5$); *kill*—weight below retrieval threshold ($w < 0.1$); *stale retention*—outdated facts that should have been invalidated but remain unchanged; *retrieval degradation*—Exact Match (EM) or F1 decrease on unrelated downstream tasks.

4.2. Scale Trend Results

Finding 2. BFS applies nonzero weight decay to 69–79% of valid memories across scales (Table 3). Hub de-

Scale	Valid	BFS Dam.	Attr Dam.	Def.	Hub _{max}
6K	344	68.9%	14.8%	78.5%	23
32K	1,744	79.0%	15.8%	80.0%	121
64K	3,414	78.5%	15.0%	80.9%	298

Table 3. Collateral damage across scales (MemoryAgentBench FactConsolidation). BFS damage increases then stabilizes; ATTR-AWARE stays at ~15%. Hub_{max}: max entity degree. Def. = 1 – Attr/BFS.

gree grows $\sim 13\times$ across scales versus $\sim 10\times$ fact growth, amplifying vulnerability. As an entity-extraction sanity check, the top hubs in these FactConsolidation runs are content entities rather than function-word artifacts: United States/America/University at 6K–32K and America/United States/University at 64K.

Finding 3. On the MemoryAgentBench scales tested here, ATTR-AWARE damage remains near ~15%. Full classification of all 51 organically damaged memories at 6K found 88.2% legitimate cascade, 11.8% borderline, and no clear false positives (see Limitations), suggesting that much of this residual reflects conservative dependency-following rather than obvious error.

4.3. Graph Topology

The degree distributions (Figure 1) are **heavy-tailed** across all scales ($\alpha \in [2.73, 2.96]$ via MLE power-law fitting (Clauset et al., 2009)), though the strict power-law hypothesis is rejected at 64K in favor of log-normal. The key driver of propagation damage is degree heterogeneity $\kappa \equiv \langle k^2 \rangle / \langle k \rangle$, which grows from 4.6 to 32.3 across scales—well above the Molloy-Reed threshold ($\kappa > 2$) that predicts catastrophic fragmentation under targeted hub removal (Molloy & Reed, 1995; Albert et al., 2000; Cohen et al., 2000).




Figure 1. Complementary cumulative degree distribution (CCDF) on log-log axes. Dotted lines: MLE power-law fits. The heavy tail extends with scale.

Graph construction control. A natural concern is whether co-occurrence clique construction inflates hub degrees. In the 6K five-fact attack setting, we compare against a *typed-relation* graph that retains only subject→object edges (no cliques): edge count decreases by 10.4%, but BFS damage drops only 1.4pp (38.4% vs 39.8%). This suggests that hub degrees are driven primarily by Zipfian entity frequency, not clique construction artifacts. Extension to 32K/64K remains for future work (see Limitations).

4.4. Cross-task Contamination

BFS propagation from high-degree named-entity hubs degrades Accurate Retrieval F1 on unrelated queries. The two highest named-entity hubs in our EventQA cross-task analysis cause −10.0pp and −11.9pp F1 drops; the strongest named hub (Debbie, degree 715) cascades through 3,381 memories. ATTR-AWARE preserves accuracy within measurement noise, confirming that directional filtering prevents cross-task interference in this evaluation. This is the most direct evidence that structural damage translates to retrieval degradation, and motivates a follow-up question: could hub-concentrated damage be deliberately triggered?

5. Hub-Targeted Amplification

5.1. Threat Model

We evaluate three attacker knowledge levels: **black-box** (general knowledge), **gray-box** (hub entity names known), and **white-box** (fact registry access). The attacker injects facts through the system’s normal memory write interface; conflict detection triggers when a new fact maps to an existing subject-key, initiating the propagation cascade under study. This write path is realistic: Mem0 and Zep expose memory-write APIs, and MINJA (Dong et al., 2025) demonstrates that even query-only interaction can induce memory writes through indirect prompt injection. We analyze this vector to inform defensive design; responsible evaluation of adversarial risks is essential for proactive safety.

5.2. Results

Attack	BFS Dam.	BFS Kill	Attr Dam.
<i>6K Context (344 valid memories)</i>			
1 fact	5.2%	0.3%	0.0%
3 facts	26.7%	4.4%	0.0%
5 facts	39.8%	17.4%	0.3%
<i>32K Context (1,744 valid memories)</i>			
1 fact	38.1%	1.0%	0.0%
5 facts	57.3%	19.3%	1.0%

Table 4. White-box adversarial injection at 6K and 32K scales. Kill: weight reduced below retrieval threshold (< 0.1). Five facts at 32K damage 57.3% of valid memories.

Finding 4. A single fake fact applies nonzero weight decay to 5.2% (6K) to 38.1% (32K) of valid memories (Table 4)— $7.3\times$ amplification from scale and hub-degree growth. Five facts at 32K affect 57.3% of valid memories, with 19.3% killed (weight < 0.1).

Attack	Facts	Hit	BFS Dam.
White-box	5	100%	39.8%
Gray-box	5	100%	20.9%
Black-box	5	0%	0.0%

Table 5. Attack comparison at 6K. Black-box fails entirely; gray-box achieves 52% of white-box damage.

Finding 5. Table 5 shows that under our random-entity black-box baseline, attacks fail to hit hub entities (0% hit rate). At minimum, gray-box knowledge of graph structure is needed for the attack vector studied here. Figure 2 shows that damage grows sublinearly with propagation depth while kill rate accelerates due to cumulative decay.




Figure 2. Propagation depth vs collateral damage (6K, 5-fact attack). Damage grows sublinearly while kill rate accelerates due to cumulative decay.

6. ATTR-AWARE Selective Propagation

6.1. Method

The core intuition is directional filtering. Consider updating “the CEO of Acme is Alice” to “... is Bob.” In this fact, *Acme* is the subject and *Alice* is the object. Facts *about Alice* (where Alice is subject)—e.g., “Alice approves Acme milestones”—may need updating because they depend on Alice’s role. Facts *about Acme* that merely co-occur with Alice—e.g., “Acme was founded in 2010”—do not, because Acme’s founding is independent of who its CEO is. ATTR-AWARE operationalizes this by following only object→subject chains: when a fact with subject *s* and object *o* is invalidated, propagation continues only to

facts where o appears as *subject* (line 4 in Algorithm 1)—restricting cascades to plausible dependency paths:

Algorithm 1 ATTR-AWARE Selective Propagation

Require: Invalidated fact f with subject s , graph G , registry R , max depth D , decay δ , current depth $h=0$

```

1:  $s \leftarrow$  subject entity of  $f$ 
2:  $E_{\text{obj}} \leftarrow$  entities linked to  $f$  in  $G$ ,  $e \neq s$ 
3: for each  $e \in E_{\text{obj}}$  do
4:    $M_{\text{down}} \leftarrow \{R[k] \mid k.\text{entity} = e\}$  {facts where  $e$  is subject}
5:   for each  $m \in M_{\text{down}}$  if  $m$  is valid do
6:      $d_{\text{in}} \leftarrow \max(\text{in\_degree}(e, G), 1)$ 
7:      $h_{\text{next}} \leftarrow h + 1$ 
8:      $m.w \leftarrow m.w \times (1 - \frac{1 - \delta^{h_{\text{next}}}}{d_{\text{in}}})$ 
9:     if  $h_{\text{next}} < D$  then
10:      Propagate( $m, G, R, D, \delta, h_{\text{next}}$ )
11:    end if
12:  end for
13: end for

```

The key filtering step is line 4: instead of visiting all graph neighbors, ATTR-AWARE retrieves only facts where the object entity appears as *subject*—restricting propagation to downstream attribute-directed chains. ATTR-AWARE does not infer causality directly; it uses attribute-keyed directionality as a conservative proxy for dependency. The in-degree adjustment (line 6) dampens decay for highly-connected entities. BFS differs by visiting all graph neighbors without semantic filtering; both methods apply in-degree dampening $(1 - (1 - \delta^{h_{\text{next}}})/d_{\text{in}})$ to reduce penalty for highly-connected entities at the next propagation hop, where $h_{\text{next}} = h + 1$. All operations are string-matching or arithmetic over pre-existing metadata—**no LLM calls** at propagation time. At 6K scale (417 nodes, 524 edges), a 5-fact attack completes propagation in 2.3 ± 0.2 ms for ATTR-AWARE versus 12.8 ± 1.2 ms for BFS—a $5.5\times$ speedup from reduced traversal scope. Propagation constitutes $<5\%$ of total processing time (dominated by LLM conflict detection), and the full pipeline processes 64K-scale graphs in <1 s without GPU.²

6.2. Propagation Strategy Comparison

We compare nine propagation strategies spanning structural, weight-based, and semantic filtering approaches (Table 6).

Finding 6. Table 6 reveals a damage spectrum from structural constraints (degree capping, modest reduction) through threshold-based filters (weighted BFS at $\theta=0.3$: 1.5% damage) to semantic filtering (ATTR-AWARE: 0.3%). Threshold

²The 64K timing is based on end-to-end pipeline observation; controlled wallclock measurements at 32K/64K scale remain for future work. The 6K micro-benchmark provides the only formally measured propagation latency.

Strategy	Dam.	Kill	EM	F1
<i>Unfiltered co-occurrence traversal</i>				
BFS (standard)	39.8%	17.4%	18.0%	18.3%
BFS + visited set	39.8%	10.5%	15.0%	15.3%
<i>Structural constraints</i>				
Degree-cap=5	24.1%	8.1%	19.0%	19.3%
Degree-cap=10	28.2%	10.8%	18.0%	18.3%
<i>Threshold-based filtering</i>				
k-hop thresh. $\theta=0.3$	18.6%	0.3%	21.0%	21.0%
k-hop thresh. $\theta=0.5$	18.6%	0.3%	21.0%	21.0%
Weighted BFS $\theta=0.1$	11.3%	3.2%	21.0%	21.3%
Weighted BFS $\theta=0.3$	1.5%	0.0%	18.0%	18.3%
<i>Semantic filtering</i>				
ATTR-AWARE	0.3%	0.0%	17.0%	17.3%

Table 6. Propagation strategy comparison (6K, 5-fact white-box attack). EM/F1 evaluated on $n=100$ adversarial queries; differences are not statistically significant (all strategies fall within 17–21%). Damage and kill metrics are deterministic.

methods require per-deployment tuning and exhibit cliff effects (k-hop $\theta=0.3$ and 0.5 produce identical results); ATTR-AWARE requires no hyperparameters. At 32K/64K scale, Weighted BFS $\theta=0.3$ maintains 0.2%/0.1% damage—comparable to ATTR-AWARE’s 1.0%/0.9%—while BFS standard damage escalates to 57.3%/61.6% (from 39.8% at 6K). Both filtered strategies remain effective even as hub degrees grow to 298; the practical advantage of ATTR-AWARE lies in hyperparameter-free operation, avoiding the cliff effects and per-deployment θ -tuning that threshold methods require. EM/F1 scores ($n=100$) converge to 17–21% across all strategies—the adversarial query set (5 injected facts, 100 questions) lacks sensitivity to discriminate what damage metrics clearly separate. The robust downstream signal is cross-task contamination (Section A): named-entity hub triggers cause -10.0pp and -11.9pp F1 drops on *unrelated* queries; ATTR-AWARE preserves accuracy within measurement noise.

Ablation: direction vs. dampening. ATTR-AWARE combines two mechanisms: (i) directional filtering (object $\rightarrow$ subject chains) and (ii) in-degree dampening $(1 - (1 - \delta^{h+1})/d_{\text{in}})$, where h is the current propagation depth and $h+1$ denotes the next hop. To isolate their contributions, we run four variants at 6K (5-fact white-box attack). Removing dampening from BFS leaves damage unchanged (39.8%) but doubles kill rate (17.4% $\rightarrow$ 33.7%); removing dampening from ATTR-AWARE has no effect (0.3% damage, 0% kill in both cases). **In this ablation, the 39.5pp damage reduction is attributable to directional filtering; dampening affects only the severity distribution, not the reach of propagation.** To test whether object $\rightarrow$ subject is the right direction, we compare three filtered variants: *downstream* (object $\rightarrow$ subject, the default), *upstream* (subject $\rightarrow$ object),

and *bidirectional*. Downstream achieves the lowest damage (0.3%, 1 memory); upstream damages 1.2% (4 memories) by propagating to other attributes of the same subject entity; bidirectional damages 1.5% (5 memories). All three tested directional filters reduce damage by $\geq 96\%$ relative to BFS, but downstream has the lowest damage in this ablation, consistent with the dependency-follows-object intuition.

6.3. Cross-experiment Consistency

Experiment	BFS	WBFS	Attr	Def.
Scale	69–79%	—	15%	78–81%
X-task F1	−10/ − 11.9pp	—	+0.0pp	100%
Adv 6K	39.8%	1.5%	0.3%	99.3%
Adv 32K	57.3%	0.2%	1.0%	98.2%
Adv 64K	61.6%	0.1%	0.9%	98.5%

Table 7. Defense effectiveness summary. **Def.** = $1 - \text{Attr}/\text{BFS}$. Scale summarizes 6K–64K FactConsolidation. Adv rows use 5 injected facts. WBFS = Weighted BFS $\theta=0.3$ (best per-scale). Cross-task: two named-entity hubs cause −10.0pp and −11.9pp F1 drops; ATTR-AWARE preserves accuracy.

Finding 7. ATTR-AWARE achieves 78–100% defense effectiveness across the reported conditions (Table 7). The strongest signal is cross-task: ATTR-AWARE preserves retrieval accuracy within measurement noise where BFS causes double-digit F1 drops under named-entity hub triggers.

6.4. Robustness to BFS Variants

A natural concern is whether BFS damage is inflated by redundant decay (a node decayed multiple times via different paths) or a particular decay parameter. We test variants at both 6K and 32K under the 5-fact white-box attack (Table 8).

Scale	BFS Variant	Dam.	Sev.	Kill
6K	Standard (decay=0.5)	39.8%	11.1%	17.4%
	+ memory visited set	39.8%	—	10.5%
	decay = 0.3	39.8%	11.9%	20.1%
	decay = 0.7	39.8%	19.8%	2.0%
	decay = 0.9	39.8%	2.0%	0.0%
32K	Standard (decay=0.5)	57.3%	—	19.3%
	+ memory visited set	57.3%	—	7.8%

Table 8. BFS damage is structurally determined: damage is identical across all variants within each scale (39.8% at 6K, 57.3% at 32K). Sev. = severe damage ($0.1 \leq w < 0.5$). At the default decay=0.5, 28.5% of valid memories suffer kill or severe damage.

Finding 8. Table 8 reveals that damage is invariant to both redundant-decay elimination and decay parameter choice at *both* scales. The 39.8% (6K) and 57.3% (32K) damage reflects BFS *reachability*—the fraction of memories structurally connected to hub entities within depth 2—not an

artifact of parameter tuning or implementation choices. At the default decay of 0.5, 28.5% of valid memories suffer severe or kill-level weight reduction ($w < 0.5$), confirming that damage is not merely cosmetic. Even at decay=0.9 (the most conservative setting), 2.0% of memories are severely damaged despite zero kills—and the 39.8% damage persists. The higher 32K rate confirms that structural vulnerability increases with scale as hub degree grows ($23 \rightarrow 121$).

7. Related Work

Memory invalidation and belief revision. Kumiho (Park, 2026) maps AGM belief revision postulates (Alchourrón et al., 1985) to a property graph with typed dependency edges, providing formal guarantees that unrelated beliefs are preserved. Our work is complementary: Kumiho shows *why* typed edges are needed (axiomatic); we show *what happens without them* (empirical), providing the damage quantification that motivates architectural choices like Kumiho’s. ATTR-AWARE’s object→subject chain is structurally analogous to justification chains in Truth Maintenance Systems (Doyle, 1979), but implemented via lightweight string-matching over pre-existing metadata rather than full dependency tracking. MaRS (Alqithami, 2025) compares six forgetting policies with provenance closure constraints on the retention side; our study addresses the invalidation side.

Memory security. MINJA (Dong et al., 2025), AgentPoison (Chen et al., 2024), and ER-MIA (Piehl et al., 2026) inject malicious content. KEPo (Chen et al., 2026) and LogicPoison (Xiao et al., 2026) manipulate graph topology. GraphRAG (Edge et al., 2024) uses community-structured graphs for retrieval; its invalidation under entity updates has not been studied. Our attack is structurally distinct: it triggers the system’s own forgetting mechanism rather than injecting content or manipulating retrieval paths.

Network resilience. Heavy-tailed networks exhibit robust-yet-fragile behavior under targeted hub removal (Albert et al., 2000; Cohen et al., 2000). Our results are consistent with this pattern in AI memory graphs ($\alpha \approx 2.7\text{--}3.0$, heavy-tailed).

8. Discussion and Conclusion

Key message. Unfiltered co-occurrence propagation is a latent risk in graph-based memory systems. It applies nonzero weight decay to 69–79% of valid memories, enables low-cost adversarial amplification under gray- or white-box hub knowledge (5 facts $\rightarrow$ 57.3% affected), and contaminates unrelated tasks under named-entity hub triggers (F1 −10.0/ − 11.9pp). ATTR-AWARE reduces scale damage by 78–81%, while both ATTR-AWARE and tuned Weighted BFS reduce adversarial damage by $\geq 98\%$; ATTR-AWARE

is hyperparameter-free, whereas Weighted BFS requires per-deployment tuning. Ablation confirms that directional filtering drives the entire damage reduction. The design principle is clear: **systems considering invalidation propagation over co-occurrence memory graphs should include dependency-aware filtering or equivalent safeguards from the outset.**

Implications. Hub entities are often essential for retrieval and cannot simply be removed, making them persistent structural vulnerabilities under unfiltered propagation. ATTR-AWARE functions as a safety guardrail that constrains propagation—including future learned policies (Alqithami, 2025)—to plausible dependency chains. Entity mention frequencies often follow heavy-tailed patterns (Zipf, 1949), producing hub-and-spoke structures wherever entities are referenced with natural frequency (Steyvers & Tenenbaum, 2005; Barabási & Albert, 1999). This makes the mechanism plausibly broader than this benchmark, but validation on additional benchmarks and production systems remains necessary.

Limitations and future work. (1) Single benchmark and single LLM: while entity extraction is rule-based and graph topology drives core findings, replication across benchmarks and model families (e.g., GPT-4o, Llama-3) is needed. This study is intended as an initial structural risk characterization; broader validation should test whether the same co-occurrence-versus-dependency failure mode appears across other memory scales, domains, and model pipelines. (2) ATTR-AWARE’s accuracy at 6K: manual inspection found 88.2% legitimate cascade, 11.8% borderline, and no clear false positives among 51 organic-damaged memories; the false negative rate (missed dependencies) remains unmeasured. The 69–79% BFS damage is an upper bound on collateral—some fraction may reflect genuine dependencies. (3) Propagation depth fixed at 2; clique construction modestly inflates hub degree (1.4pp effect at 6K); both warrant multi-scale analysis. (4) Our experiments are text-memory only; extending the same propagation analysis to multimodal memory graphs with visual, spatial, or embodied entities is an important next step. (5) Our direction ablation confirms object→subject is optimal among the three structural variants tested; embedding-similarity-based filters and validation on production systems (Mem0 graph mode, Zep) once they adopt propagation are important next steps.

LLM/Agent Usage Disclosure

This research used LLM-based tools in three capacities: (1) Literature search and survey assistance (paper discovery and classification); (2) LaTeX editing and proofreading; (3) Experimental code development (benchmark integration,

graph analysis scripts). All experimental results, analysis, and conclusions were produced and validated by the author. The core experimental pipeline (entity extraction, subject-key matching, propagation algorithms) uses rule-based methods without LLM involvement at runtime, except for conflict detection which uses Gemma-4-31B-it as specified in the experimental setup.

References

- Albert, R., Jeong, H., and Barabási, A.-L. Error and attack tolerance of complex networks. *Nature*, 406:378–382, 2000.
- Alchourrón, C. E., Gärdenfors, P., and Makinson, D. On the logic of theory change: Partial meet contraction and revision functions. *Journal of Symbolic Logic*, 50(2): 510–530, 1985. doi: 10.2307/2274239.
- Alqithami, S. Forgetful but faithful: MaRS – a memory-aware retention schema for AI agents. *arXiv preprint arXiv:2512.12856*, 2025.
- Barabási, A.-L. and Albert, R. Emergence of scaling in random networks. *Science*, 286(5439):509–512, 1999.
- Chen, A., Yang, C., Kongsomsaksakul, S., and Lee, M. Network-based accessibility measures for vulnerability analysis of degradable transportation networks. *Networks and Spatial Economics*, 7(3):241–256, 2007.
- Chen, Q., Qi, C., Huang, Y., Li, M., Wang, R., Zhang, D., Qin, K., and Liang, S. KEPO: Knowledge evolution poison on GraphRAG. In *Proceedings of the ACM Web Conference (WWW)*, 2026.
- Chen, Z., Xiang, Z., Xiao, C., Song, D., and Li, B. Agent-Poison: Red-teaming LLM agents via poisoning memory or knowledge bases. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- Chhikara, P., Khant, D., Aryan, S., Singh, T., and Yadav, D. Mem0: Building production-ready AI agents with scalable long-term memory. In *Proceedings of ECAI 2025*, pp. 2993–3000, 2025.
- Clauset, A., Shalizi, C. R., and Newman, M. E. Power-law distributions in empirical data. *SIAM Review*, 51(4): 661–703, 2009.
- Cohen, R., Erez, K., ben Avraham, D., and Havlin, S. Resilience of the internet to random breakdowns. *Physical Review Letters*, 85(21):4626–4628, 2000.
- Devarangadi Sunil, B., Sinha, I., Maheshwari, P., Todmal, S., Mallik, S., and Mishra, S. Memory poisoning attack and defense on memory based LLM-agents. *arXiv preprint arXiv:2601.05504*, 2026.

- Dong, S., Xu, S., He, P., Li, Y., Tang, J., Liu, T., Liu, H., and Xiang, Z. MINJA: Memory INjection attack against LLM agents via query-only interaction. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- Doyle, J. A truth maintenance system. *Artificial Intelligence*, 12(3):231–272, 1979.
- Edge, D., Trinh, H., Cheng, N., Bradley, J., Chao, A., Mody, A., Truitt, S., and Larson, J. From local to global: A graph RAG approach to query-focused summarization. *arXiv preprint arXiv:2404.16130*, 2024.
- He, Y. et al. MemoryArena: Benchmarking agent memory in interdependent multi-session agentic tasks. *arXiv preprint arXiv:2602.16313*, 2026.
- Hu, Y., Wang, Y., and McAuley, J. Evaluating memory in LLM agents via incremental multi-turn interactions. In *International Conference on Learning Representations (ICLR)*, 2026.
- Jiang, D., Li, Y., Li, G., and Li, B. MAGMA: A multi-graph based agentic memory architecture for AI agents. *arXiv preprint arXiv:2601.03236*, 2026.
- MITRE Corporation. MITRE ATLAS: Adversarial threat landscape for AI systems. <https://atlas.mitre.org>, 2023.
- Molloy, M. and Reed, B. A critical point for random graphs with a given degree sequence. *Random Structures & Algorithms*, 6(2-3):161–180, 1995.
- Park, Y. B. Graph-native cognitive memory for AI agents: Formal belief revision semantics for versioned memory architectures. *arXiv preprint arXiv:2603.17244*, 2026.
- Piehl, M., Xi, Z., Xiong, Z., He, P., and Ye, M. ER-MIA: Black-box adversarial memory injection attacks on long-term memory-augmented LLMs. *arXiv preprint arXiv:2602.15344*, 2026.
- Rasmussen, P., Paliychuk, P., Beauvais, T., Ryan, J., and Chalef, D. Zep: A temporal knowledge graph architecture for agent memory. *arXiv preprint arXiv:2501.13956*, 2025.
- Steyvers, M. and Tenenbaum, J. B. The large-scale structure of semantic networks: Statistical analyses and a model of semantic growth. *Cognitive Science*, 29(1):41–78, 2005. doi: 10.1207/s15516709cog2901_3.
- Xiao, Y., Chen, J., Zhang, Q., Zhang, Y., Zhou, C., Yang, L., Ren, L., Yang, X., and Huang, X. LogicPoison: Logical attacks on graph retrieval-augmented generation. *arXiv preprint arXiv:2604.02954*, 2026.
- Xiong, Z., Lin, Y., Xie, W., He, P., Liu, Z., Tang, J., Lakkaraju, H., and Xiang, Z. How memory management impacts LLM agents: An empirical study of experience-following behavior. *arXiv preprint arXiv:2505.16067*, 2025.
- Xu, W., Liang, Z., Mei, K., Li, H., Luo, H., and Chen, X. Amem: Agentic memory for LLM agents. *arXiv preprint arXiv:2502.12110*, 2025.
- Xu, Z., Zhu, X., Yao, Y., Xue, M., and Song, Y. From storage to steering: Memory control flow attacks on LLM agents. *arXiv preprint arXiv:2603.15125*, 2026.
- Yang, C., Zhou, C., Xiao, Y., et al. Graph-based agent memory: Taxonomy, techniques, and applications. *arXiv preprint arXiv:2602.05665*, 2026.
- Zipf, G. K. *Human Behavior and the Principle of Least Effort*. Addison-Wesley, 1949.

A. Cross-task Contamination (C2)

We memorize 64K EventQA turns, then trigger propagation from high-degree hubs in isolation (snapshot-restore between hubs) and evaluate Accurate.Retrieval queries ($n=100$) that are *unrelated* to the propagation trigger. In this diagnostic, the degree ranking surfaces three sentence-initial function-word or pronoun artifacts due to conservative rule-based extraction. We therefore report an artifact-aware named-entity subset below, which still shows substantial cross-task contamination. This sanity check is not a substitute for full stopword/pronoun-filtered multi-scale replication; it indicates that cross-task contamination is not solely driven by function-word hubs. All main conclusions rely on reachability and propagation comparisons, not on the identity of any single raw hub. Baseline F1 = 70.9%.

Hub	Deg.	Aff.	BFS F1	$\Delta F1$	Attr F1
Debbie	715	3,381	59.0%	−11.9pp	70.9%
Marianne	589	3,064	61.0%	−10.0pp	73.0%

Table 9. Named-entity hub cross-task contamination (EventQA 64K $\rightarrow$ Accurate.Retrieval). **Aff.**: memories affected by BFS propagation from each hub. **$\Delta F1$** : BFS F1 minus baseline. ATTR-AWARE preserves accuracy within measurement noise on this $n=100$ evaluation.

Across these two named hubs, BFS causes −10.0pp and −11.9pp F1 degradation on unrelated queries, with the worst hub (Debbie, degree 715) cascading to 3,381 memories. ATTR-AWARE preserves accuracy within measurement noise: F1 remains close to the 70.9% baseline.

B. Multi-run Statistics

We run the 6K adversarial attack 7 times (19 paired observations). Pooled Wilcoxon signed-rank test: $W = 86.0$, $p = 0.0022$; bootstrap 95% BCa CI [2.82, 15.67]pp excludes zero. A mixed-effects model with attack strength as fixed effect and run as random intercept confirms the 5-fact condition is significant ($p = 0.020$).

C. Graph Construction

Entity extraction uses rule-based proper noun detection and “the X of Y” patterns. Co-occurrence edges form cliques per fact. Subject-keys are deterministic: “the president of Russia is Putin” $\rightarrow$ Russia:president. Conflict is triggered when a new fact maps to an existing subject-key.

D. Experimental Configuration

LLM: Gemma-4-31B-it (temp=0.0). Embedding: BAAI/bge-m3 (1024-dim). Graph: NetworkX DiGraph. Storage: SQLite in-memory. Benchmark: MemoryAgent-Bench (ICLR 2026). Scales: 6K/32K/64K. Propagation: depth=2, decay=0.5.